

# Instrumenting Applications with Metrics for Prometheus

---

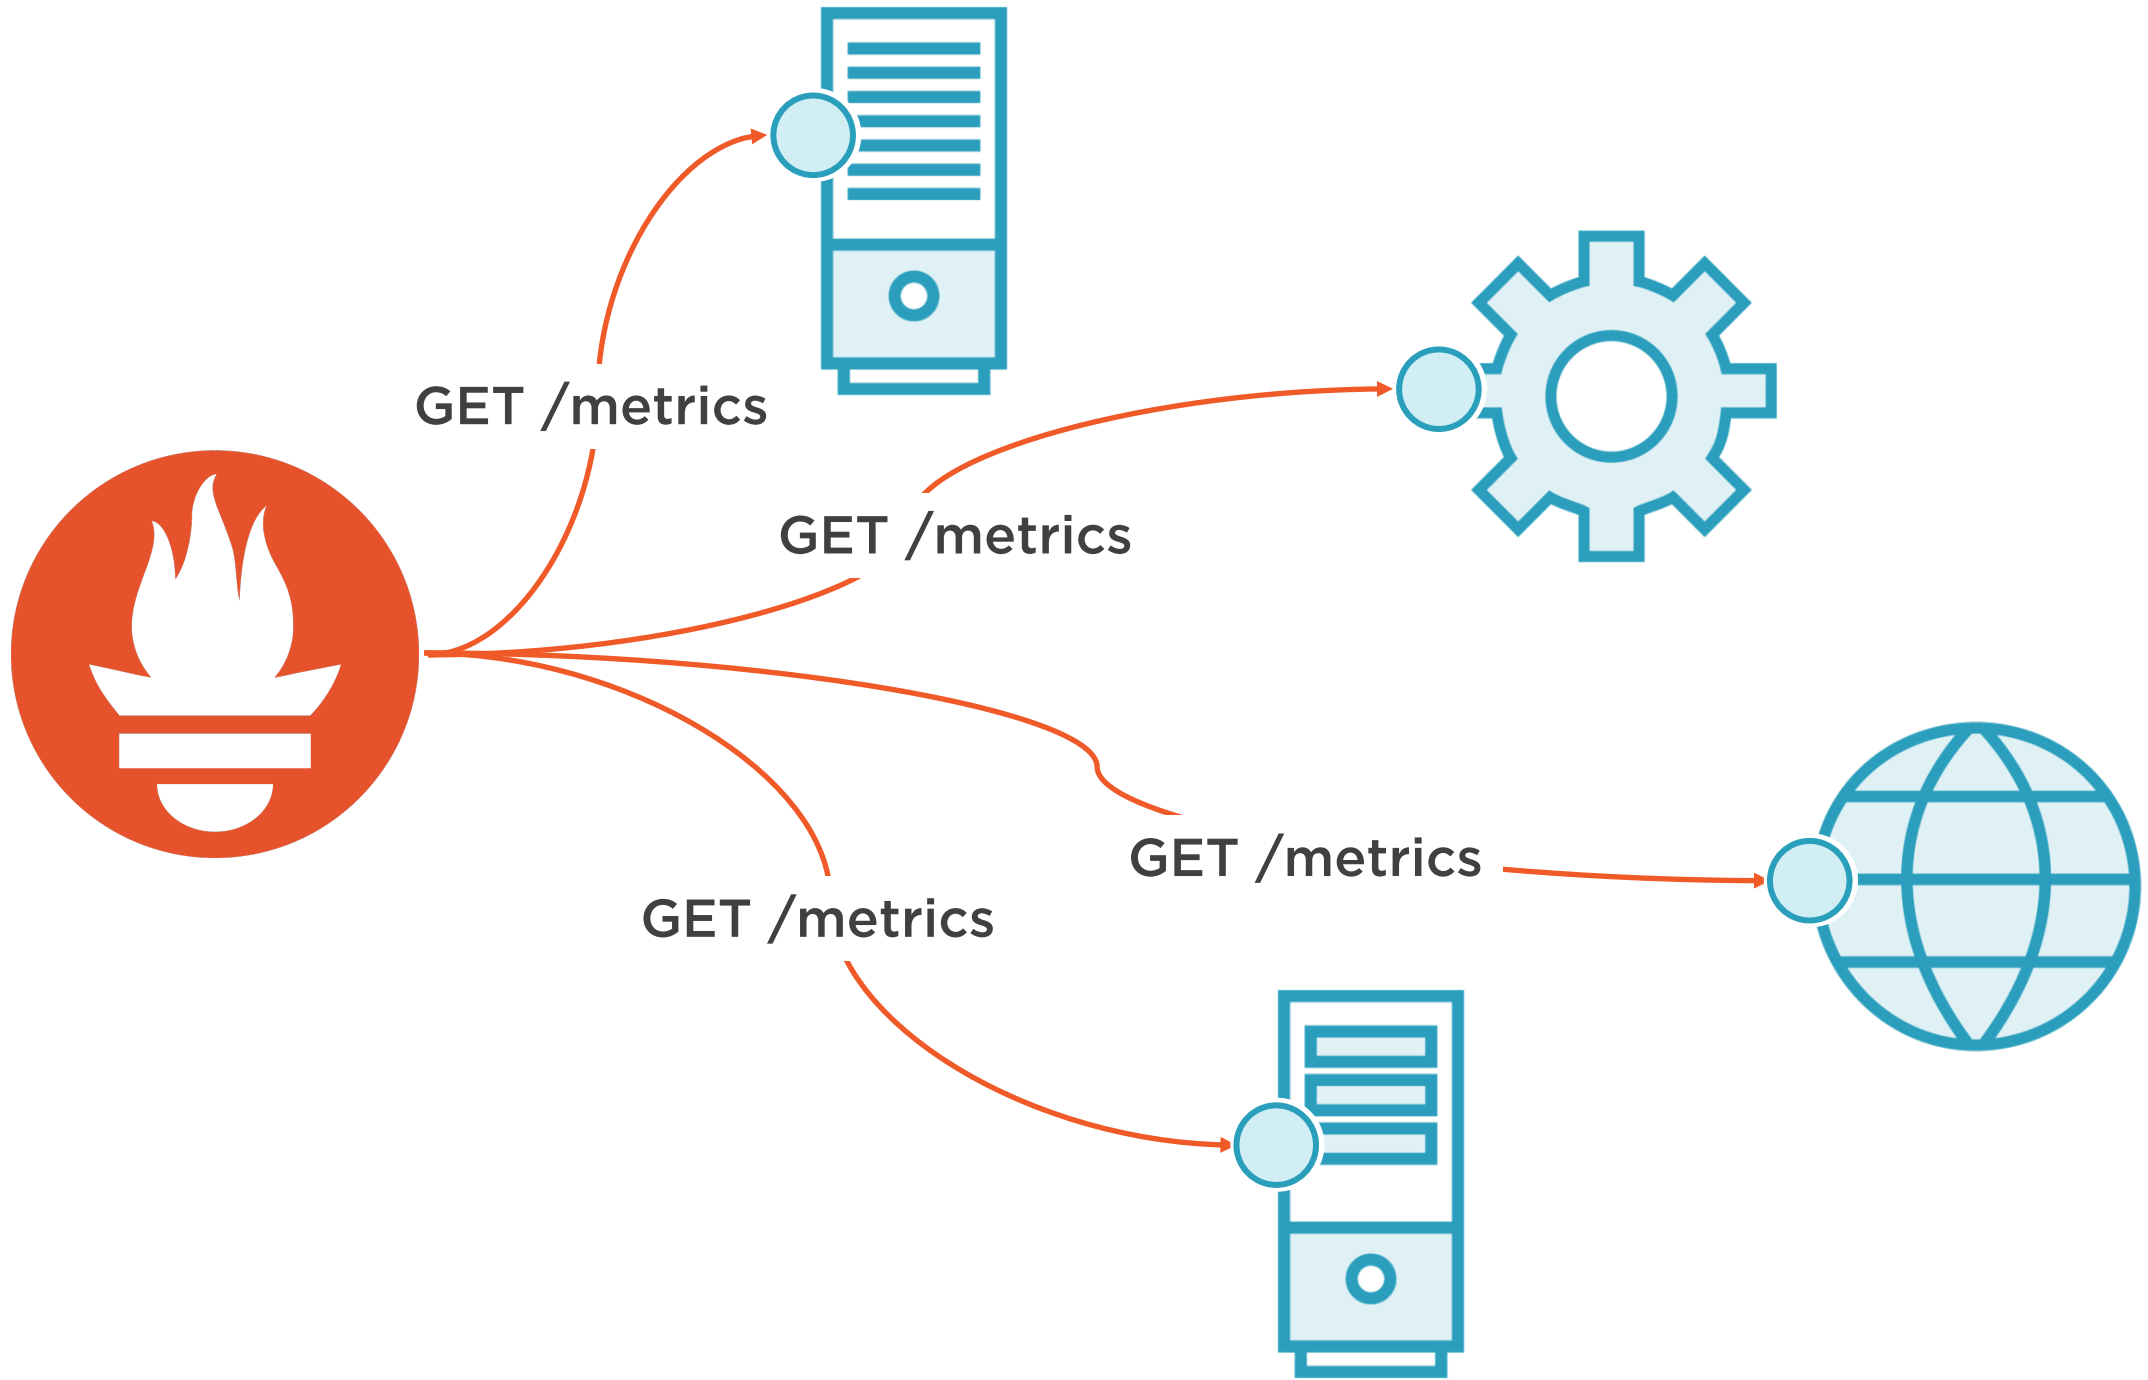
ADDING INSTRUMENTATION WITH CLIENT LIBRARIES

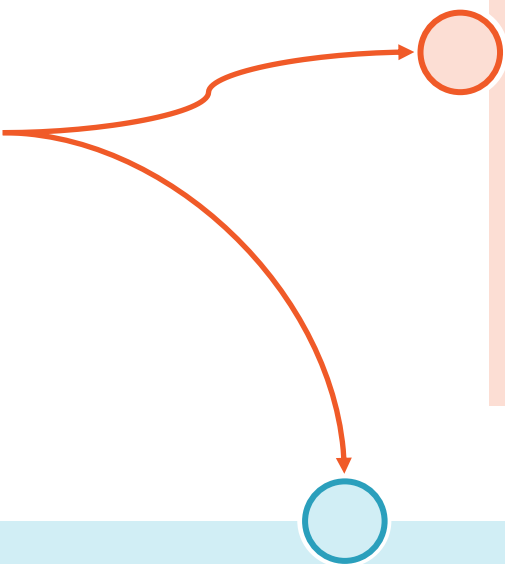


**Elton Stoneman**

CONSULTANT & TRAINER

@EltonStoneman | [blog.sixeyed.com](http://blog.sixeyed.com)

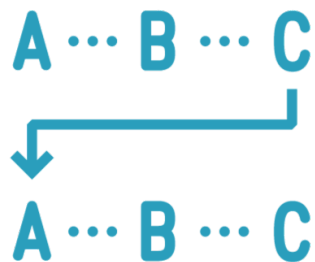




## Client libraries



## Exporters



# Getting Started with Prometheus

★★★★★ By Elton Stoneman

Prometheus is the preferred monitoring tool for containers, but it works just as well in any environment. This course will teach you how to get up and running with Prometheus and add a consistent monitoring approach to all your apps and servers.

```
tar xvfz node_exporter-1.0.0.linux-  
amd64.tar.gz
```

```
cd node_exporter-1.0.0.linux-amd64/
```

Run (normally this would be a daemon):

```
./node_exporter
```

Browse to <http://ns-prom-ub18049100/metrics>

```
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=textfile  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=thermal_zone  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=time  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=ratelimit  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=udp_queues  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=uname  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=vmstat  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=xfs  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:112 collector=zfs  
level=info ts=2020-06-05T09:43:35.794Z caller=node_export  
er.go:191 msg="listening on" address=:9100  
level=info ts=2020-06-05T09:43:35.794Z caller=tls_config.  
go:170 msg="TLS is disabled and it cannot be enabled on t  
he fly." http2=false
```

## Course info

Rating ★★★★★ (19)

Level Beginner

Updated Jun 24, 2020

Duration 1h 49m

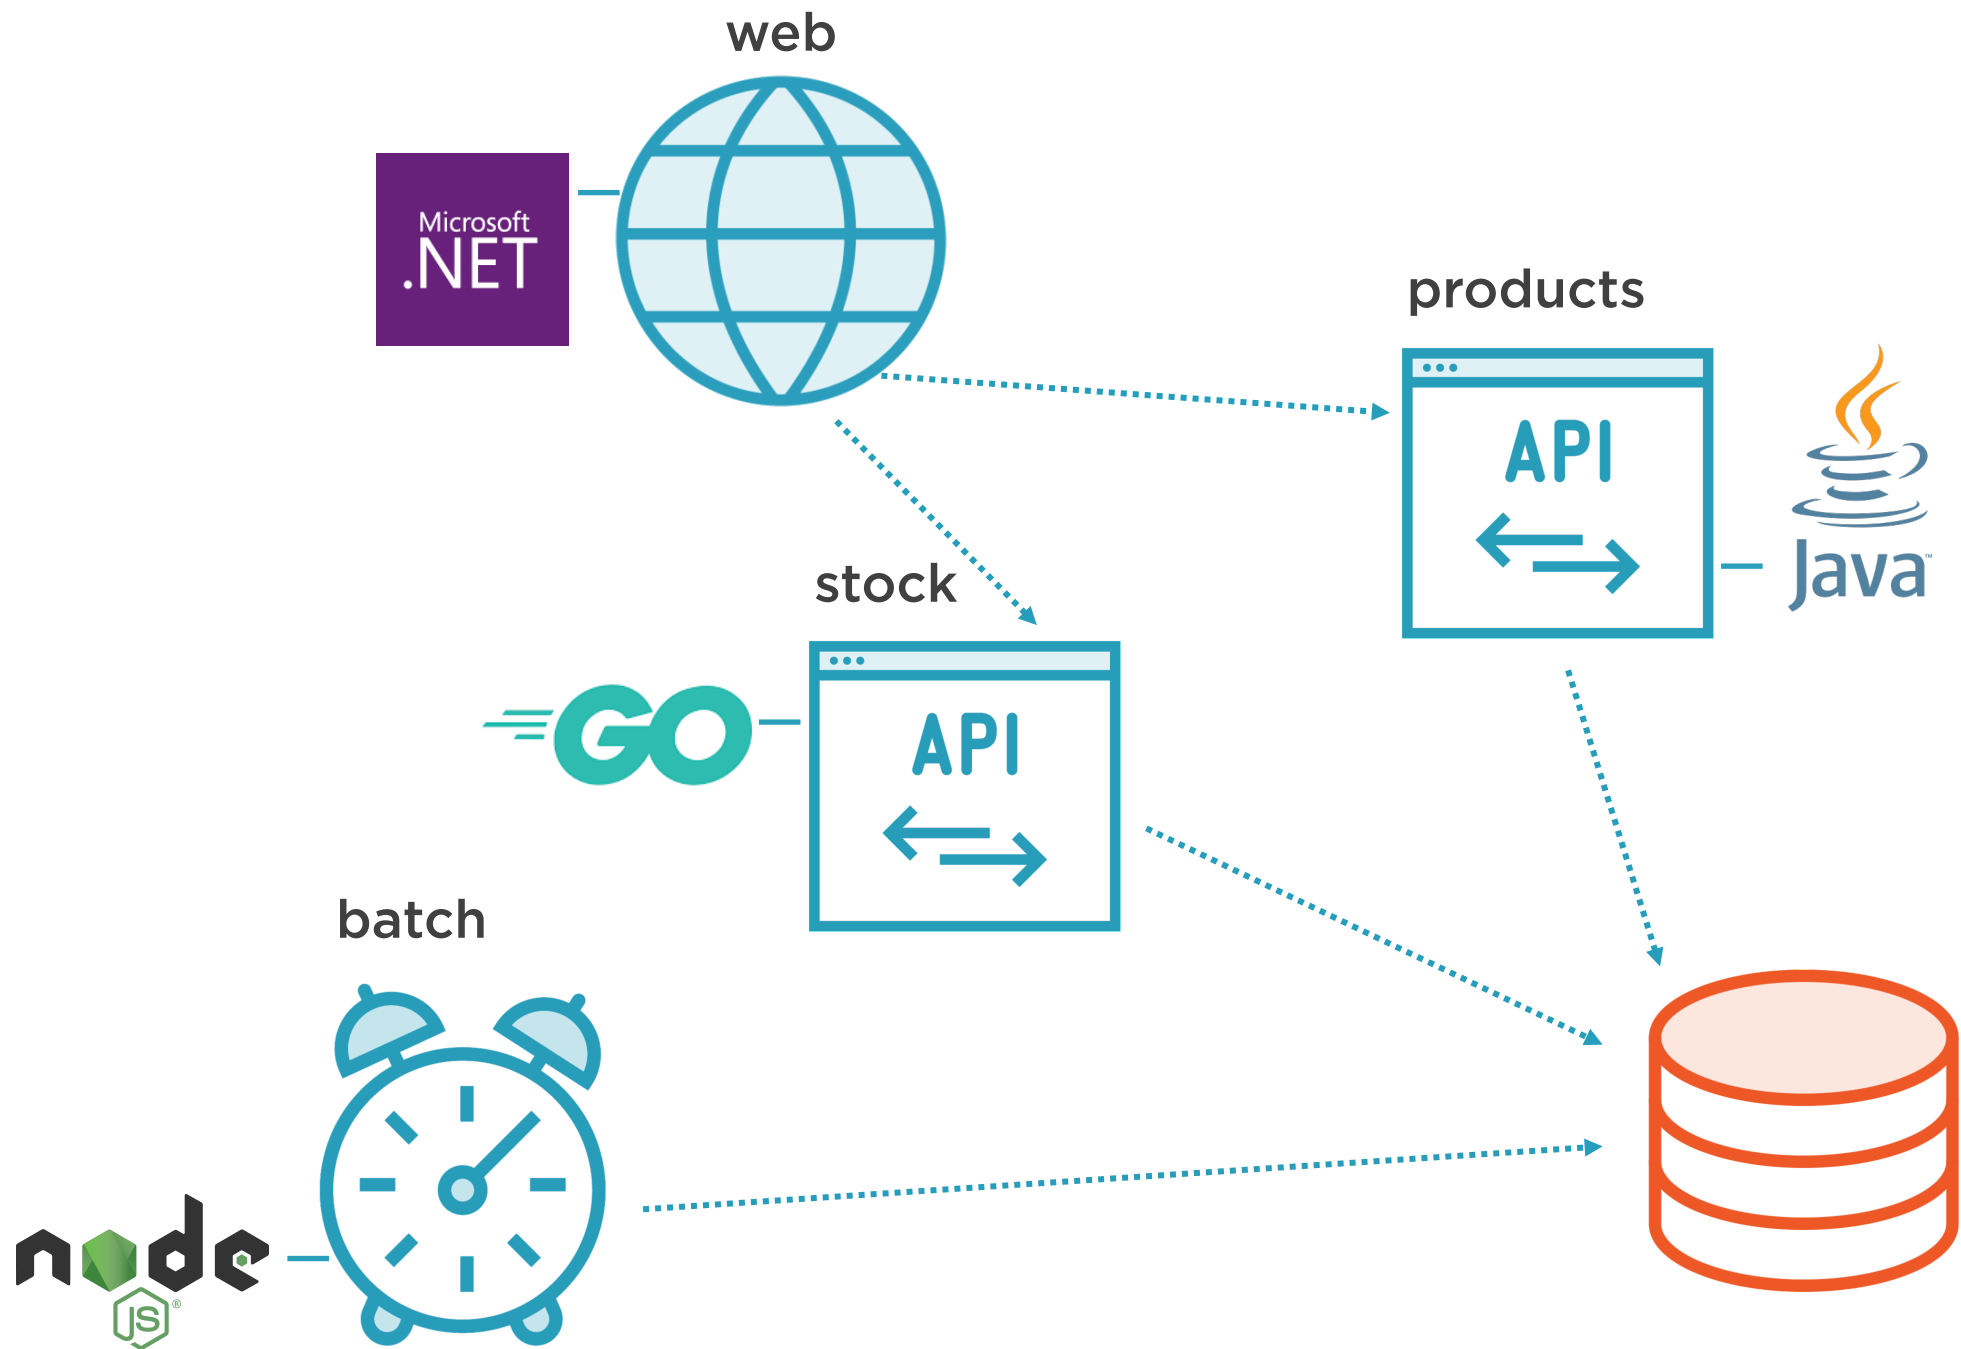
## Description

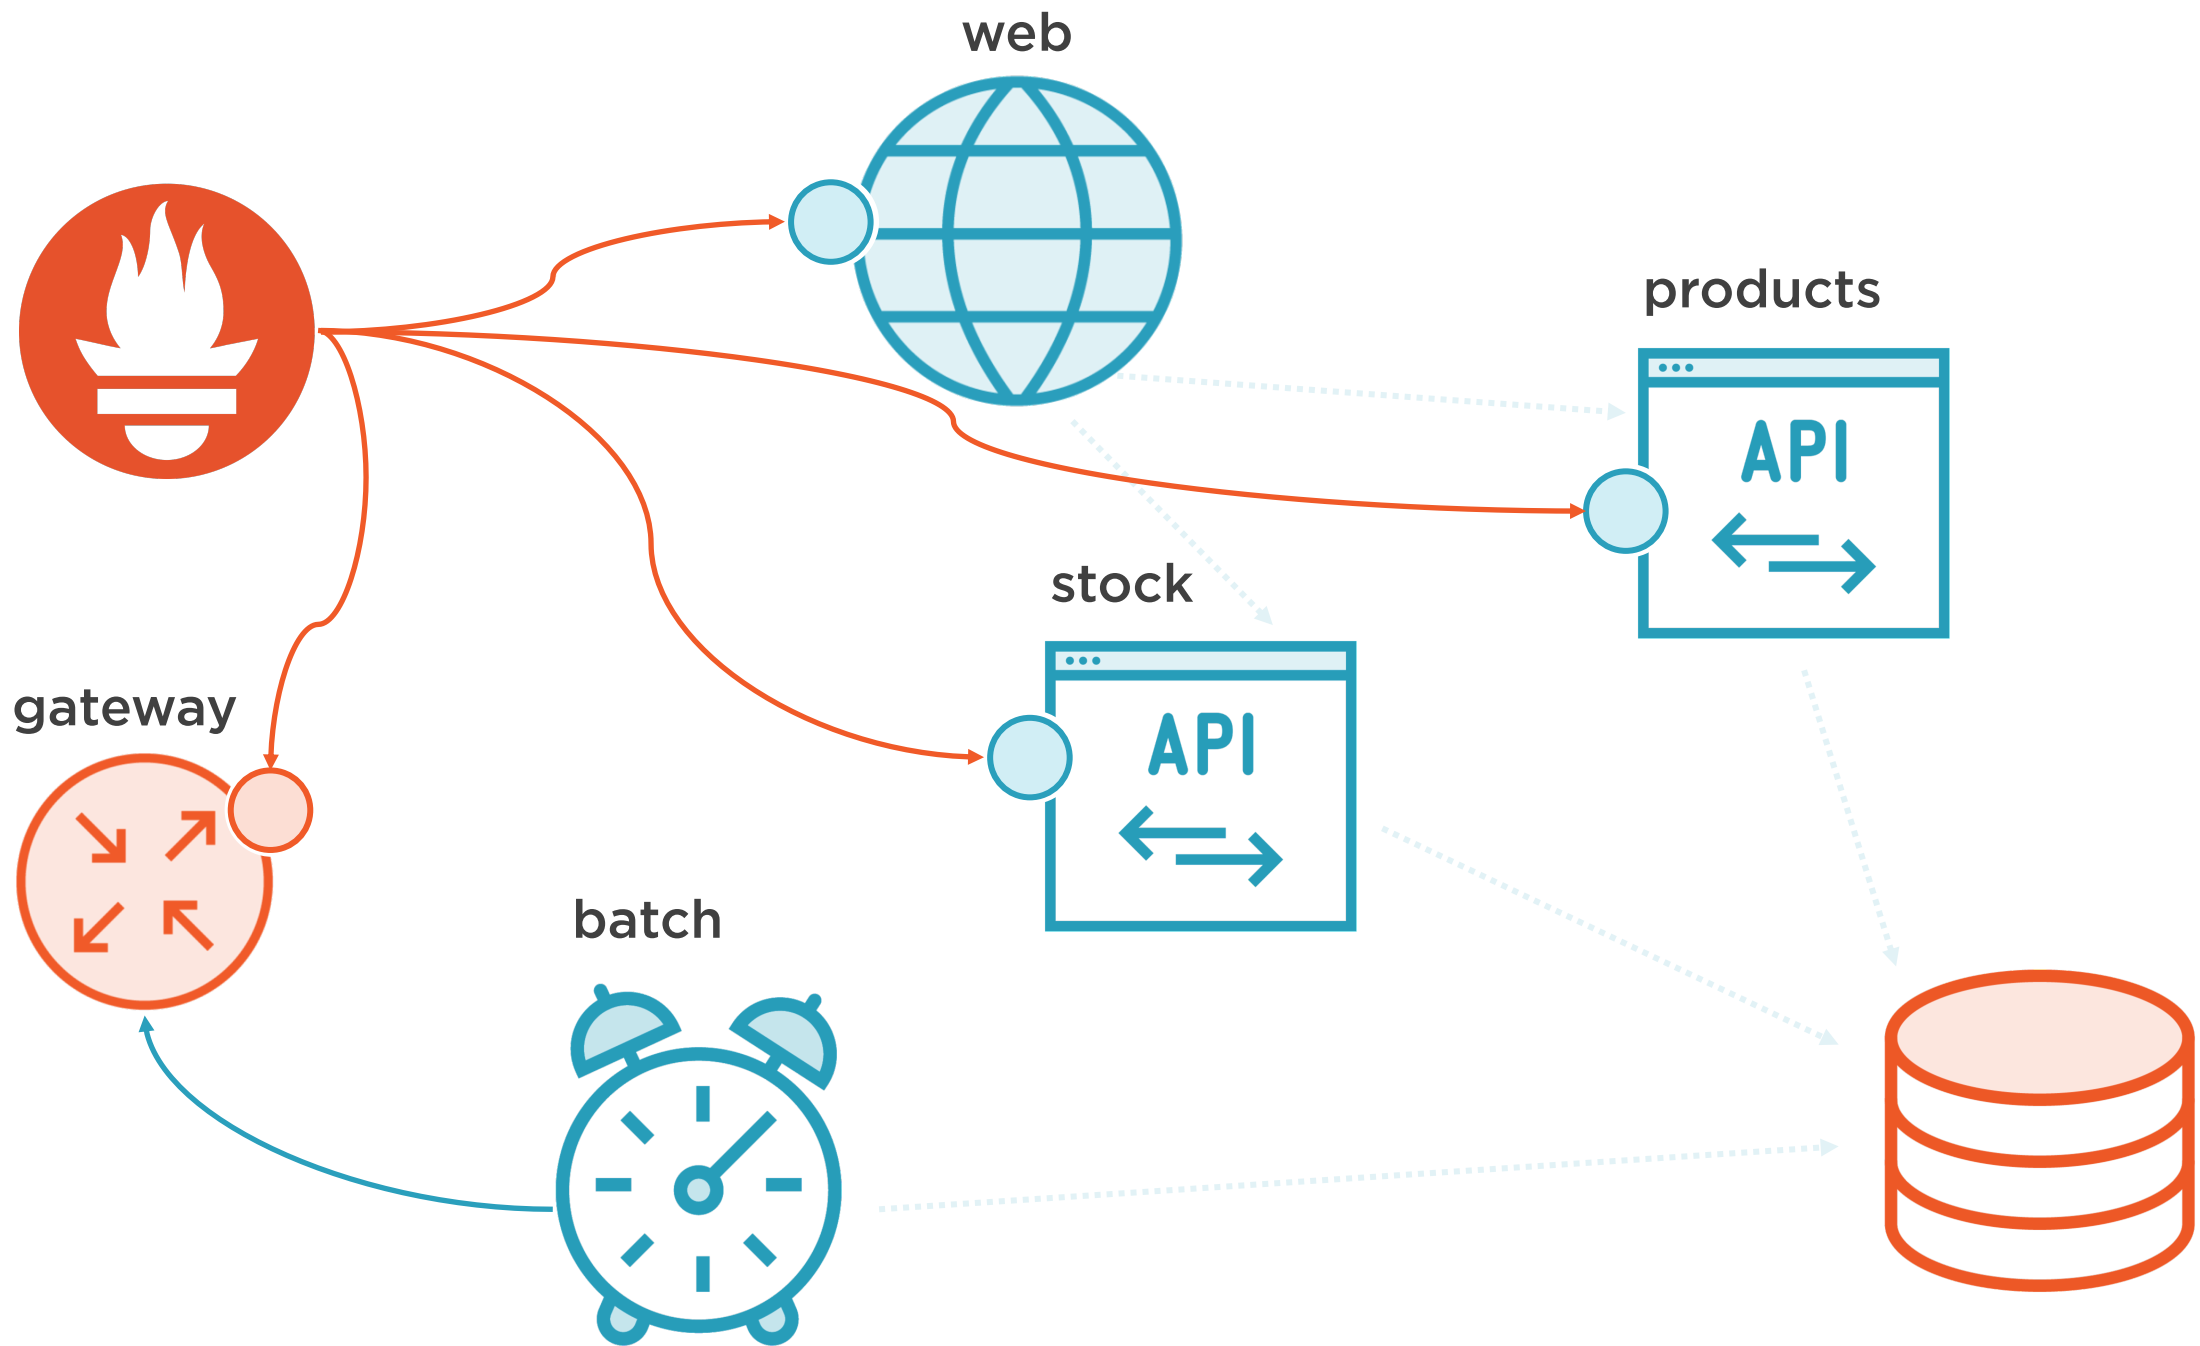
Prometheus is a cross-platform monitoring tool that lets you collect metrics from servers, containers, and applications and work with them all in the same way. In this course, Getting Started with Prometheus, you'll learn why it's such a popular approach to monitoring and how you can start bringing it into your organization. First, you'll learn about the architecture of Prometheus and how it uses a pull model to collect metrics from many targets. Then, you'll explore how to produce metrics from Linux and Windows servers using an exporter utility and from applications using a client library, and how to configure Prometheus to fetch those metrics. Finally, you'll discover the query language PromQL, how you can use it to track the changes in metrics over time, and visualize all the metrics in a dashboard. When you've finished with the course, you'll have the basic skills and knowledge of Prometheus needed to run a trial and evaluate it for your organization.

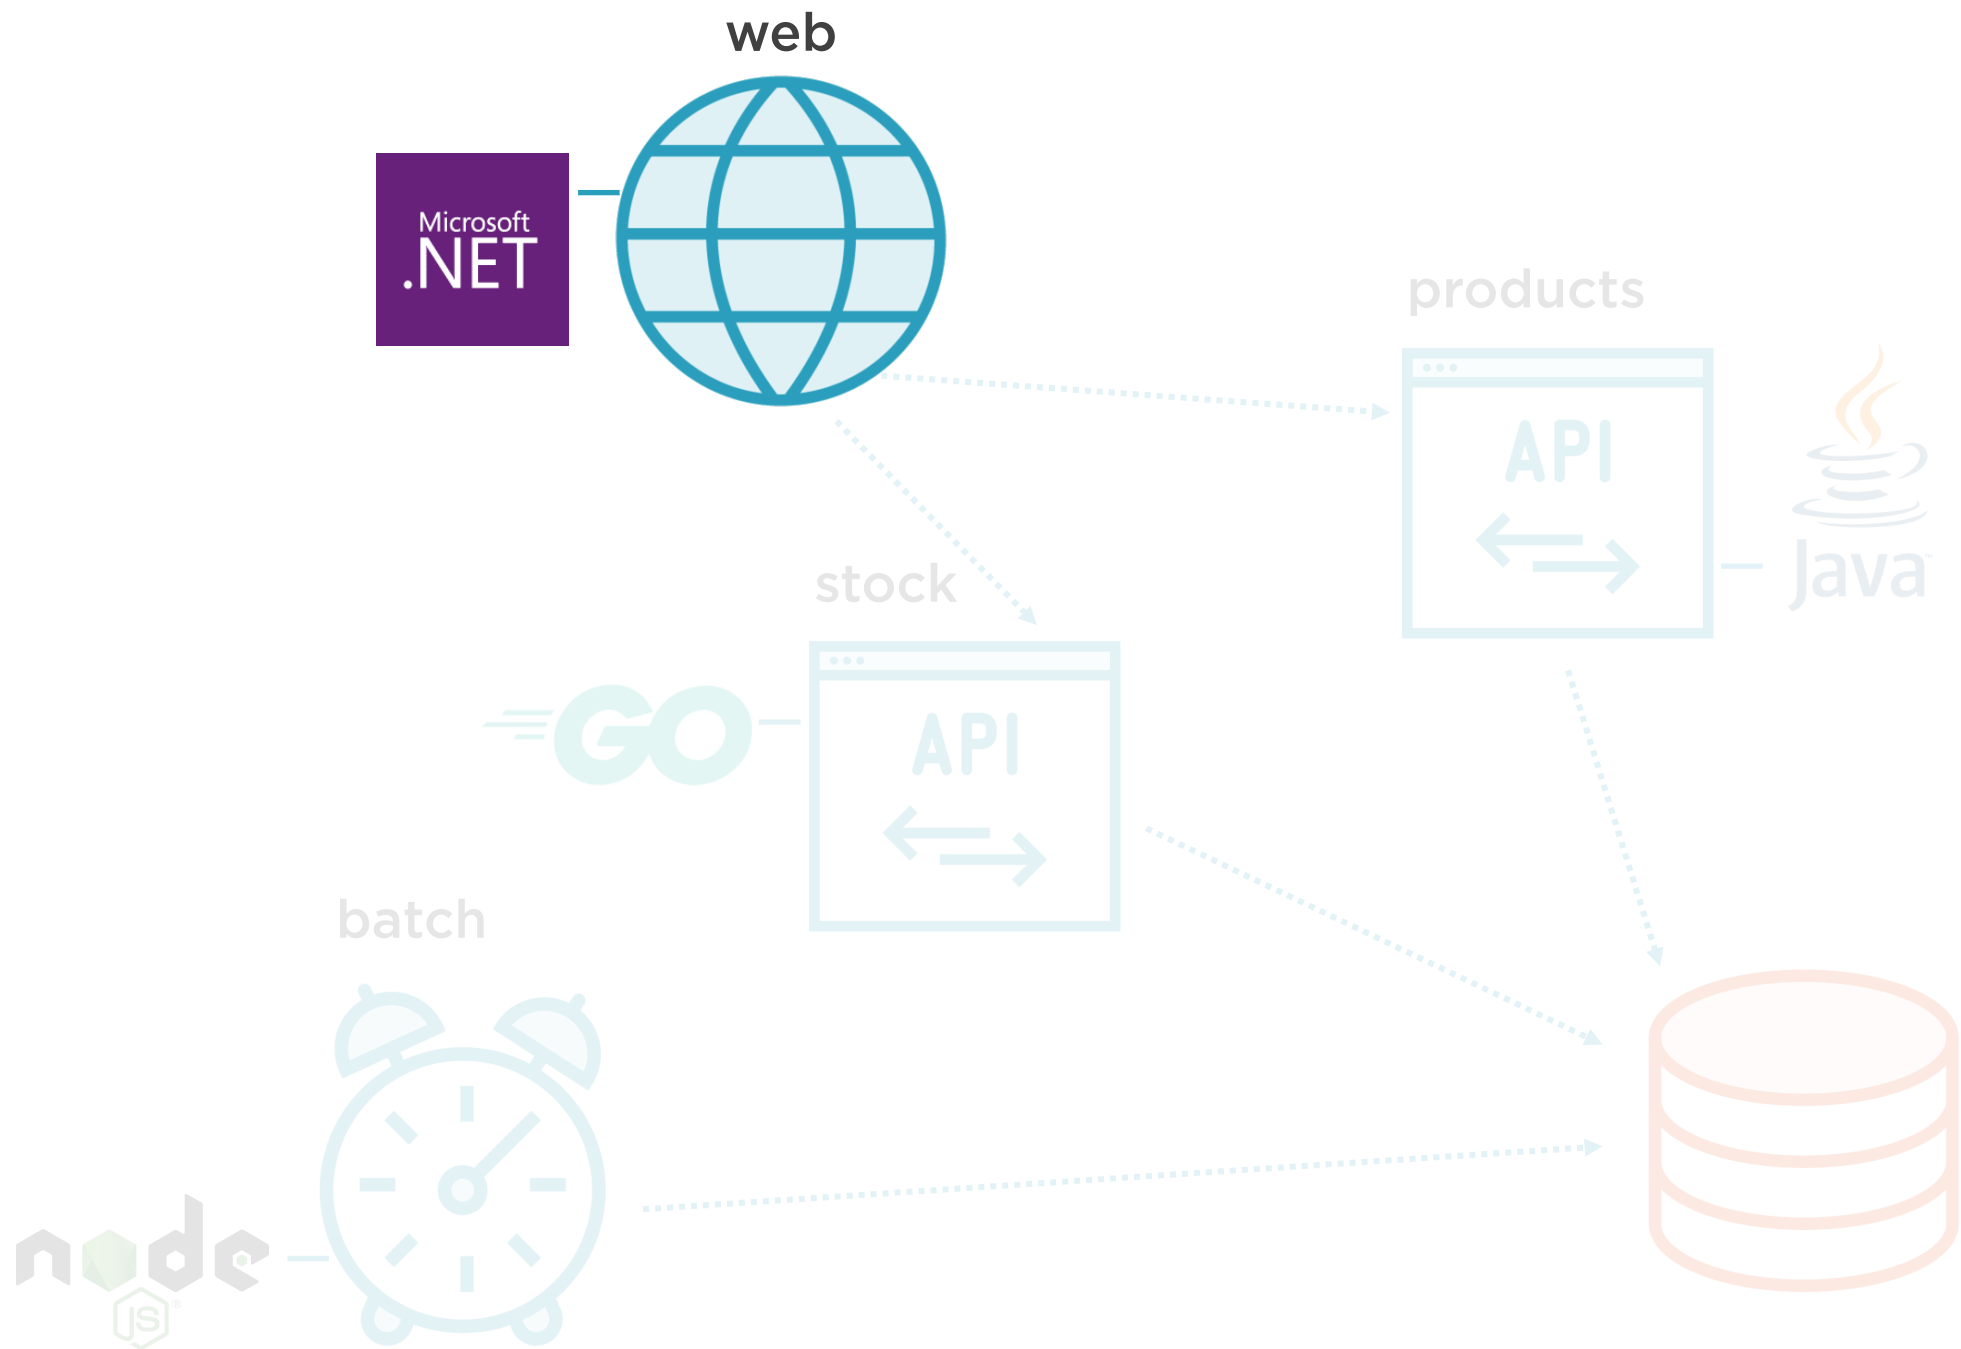
## Understanding How Prometheus Works

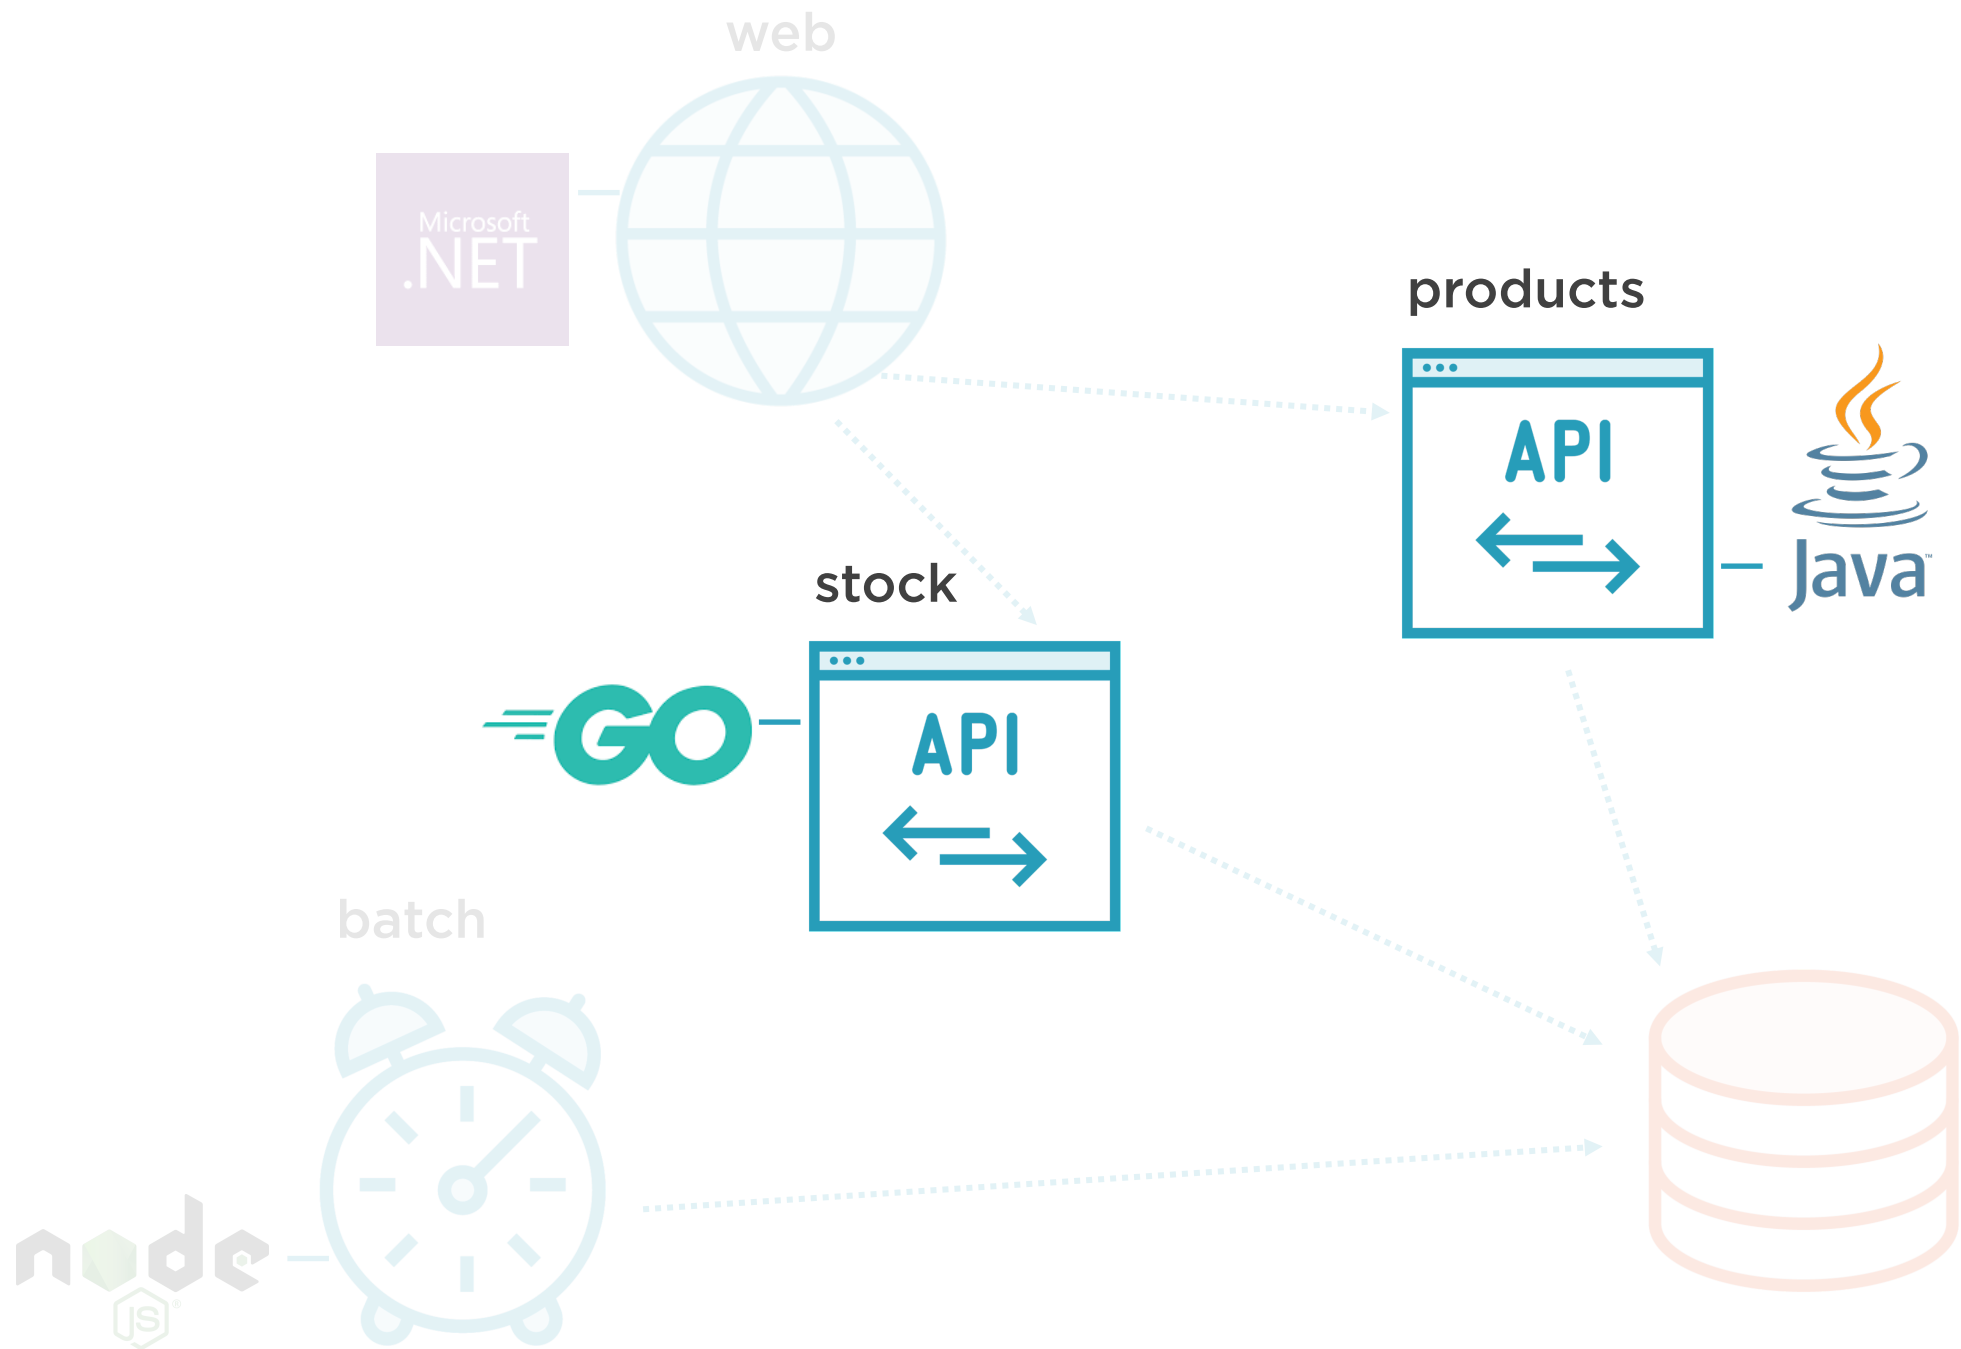
- Demo: Counters and Gauges in Linux Apps 5m
- Demo: Summaries and Histograms in Windows Apps 7m
- Exploring the Prometheus Architecture 3m
- Module Summary 2m
- Understanding Labels and Data Granularity 5m
- Understanding the Prometheus Metric Types 6m
- What Makes Prometheus so Awesome? 4m

<https://is.gd/yitezi>

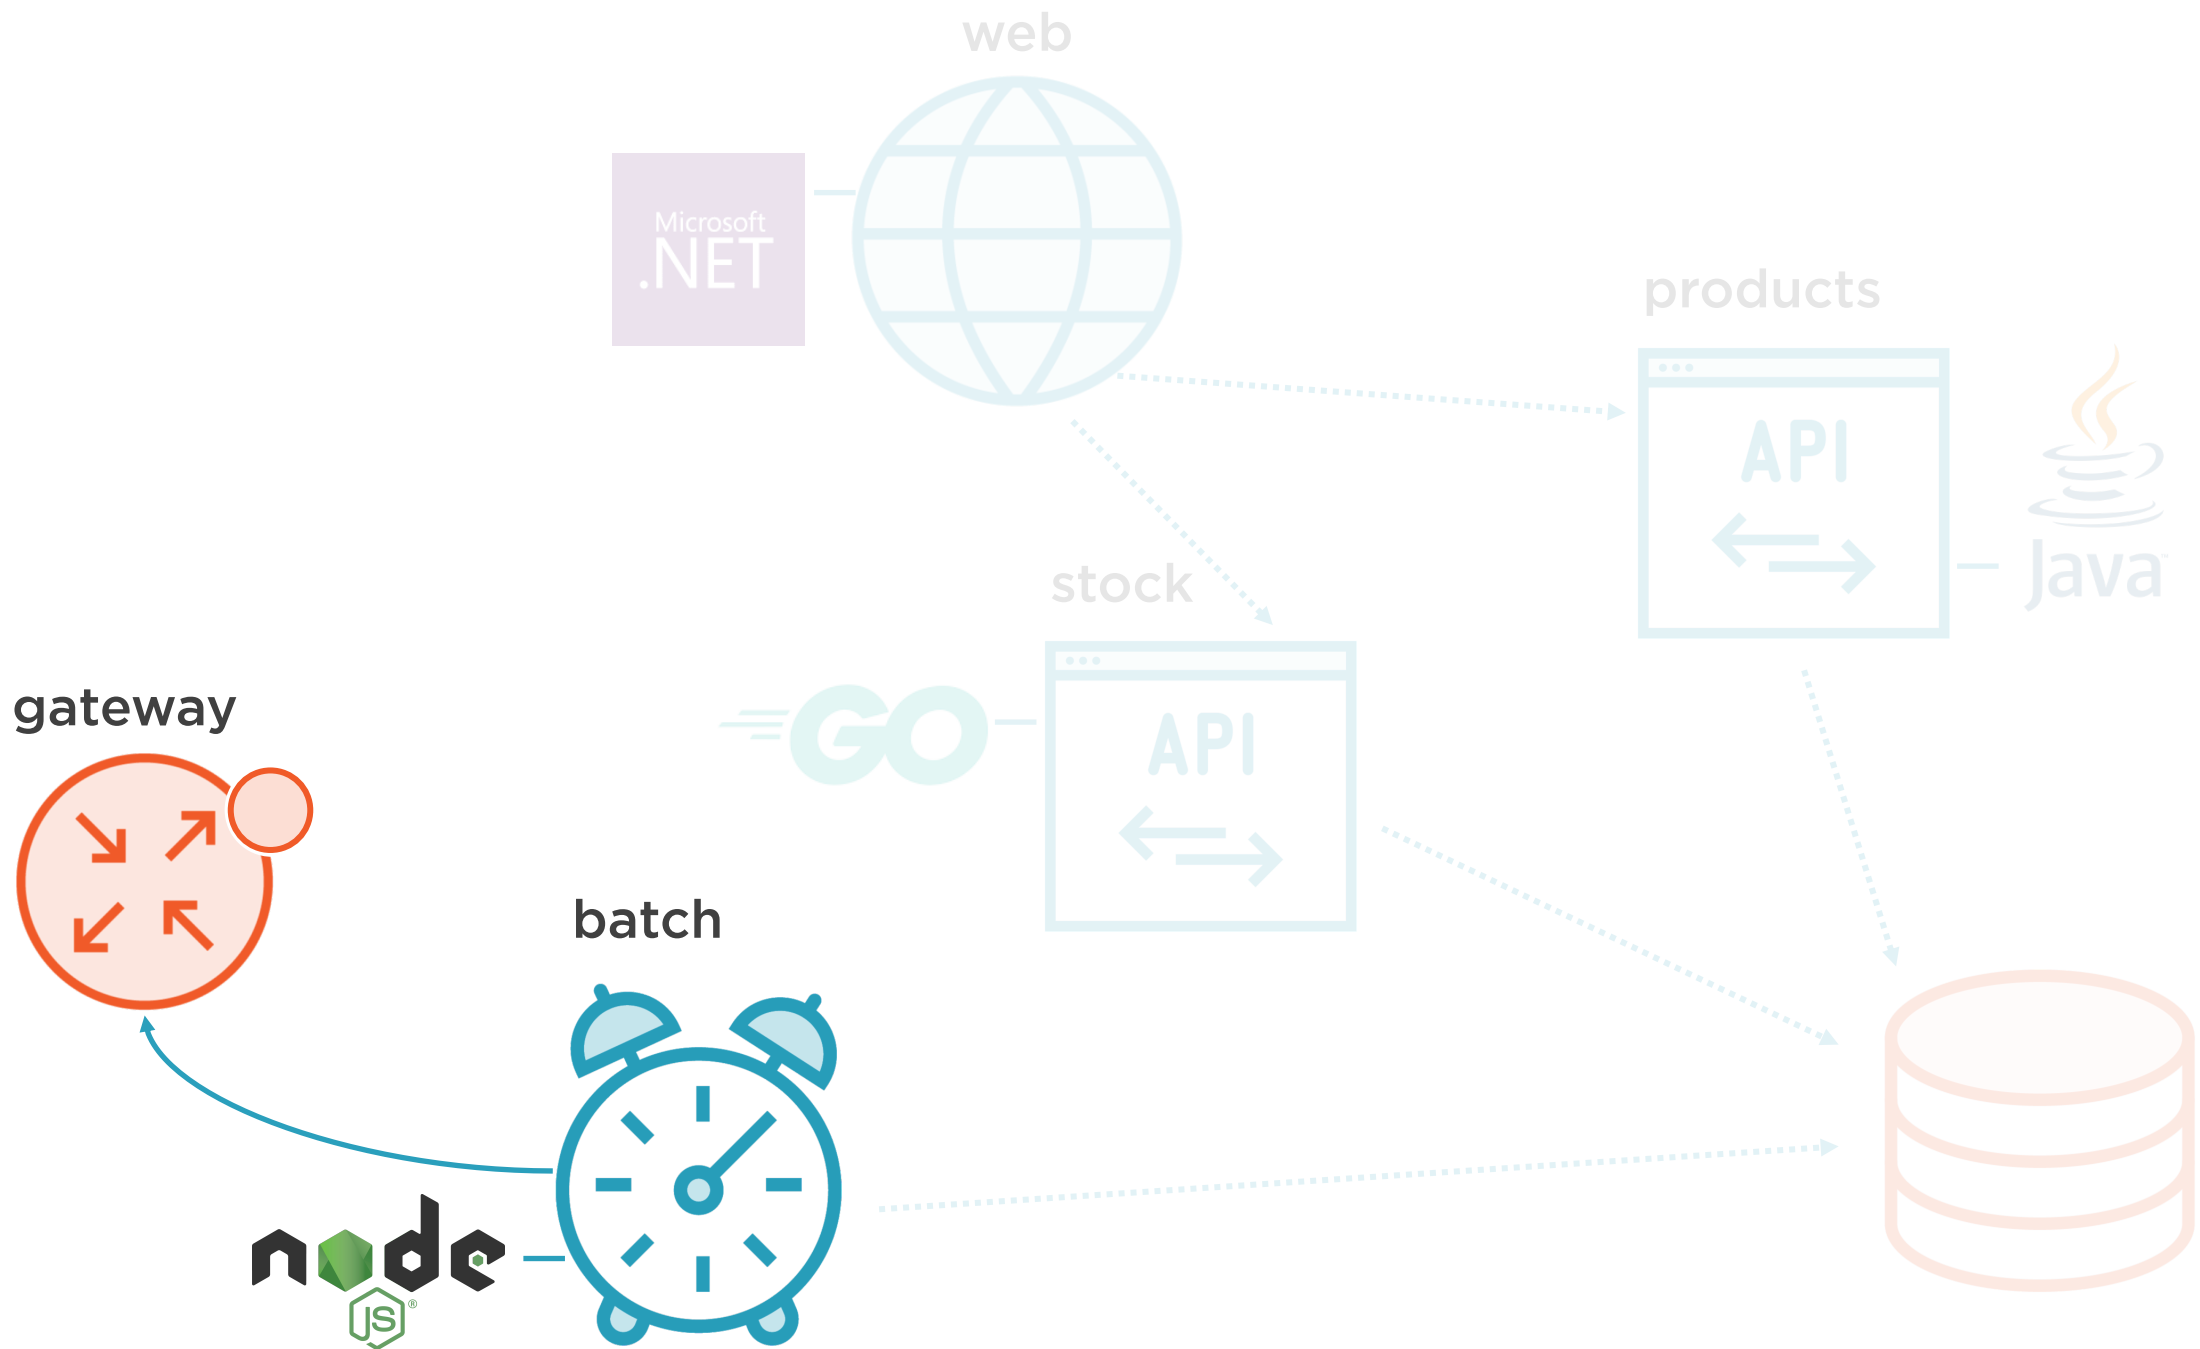


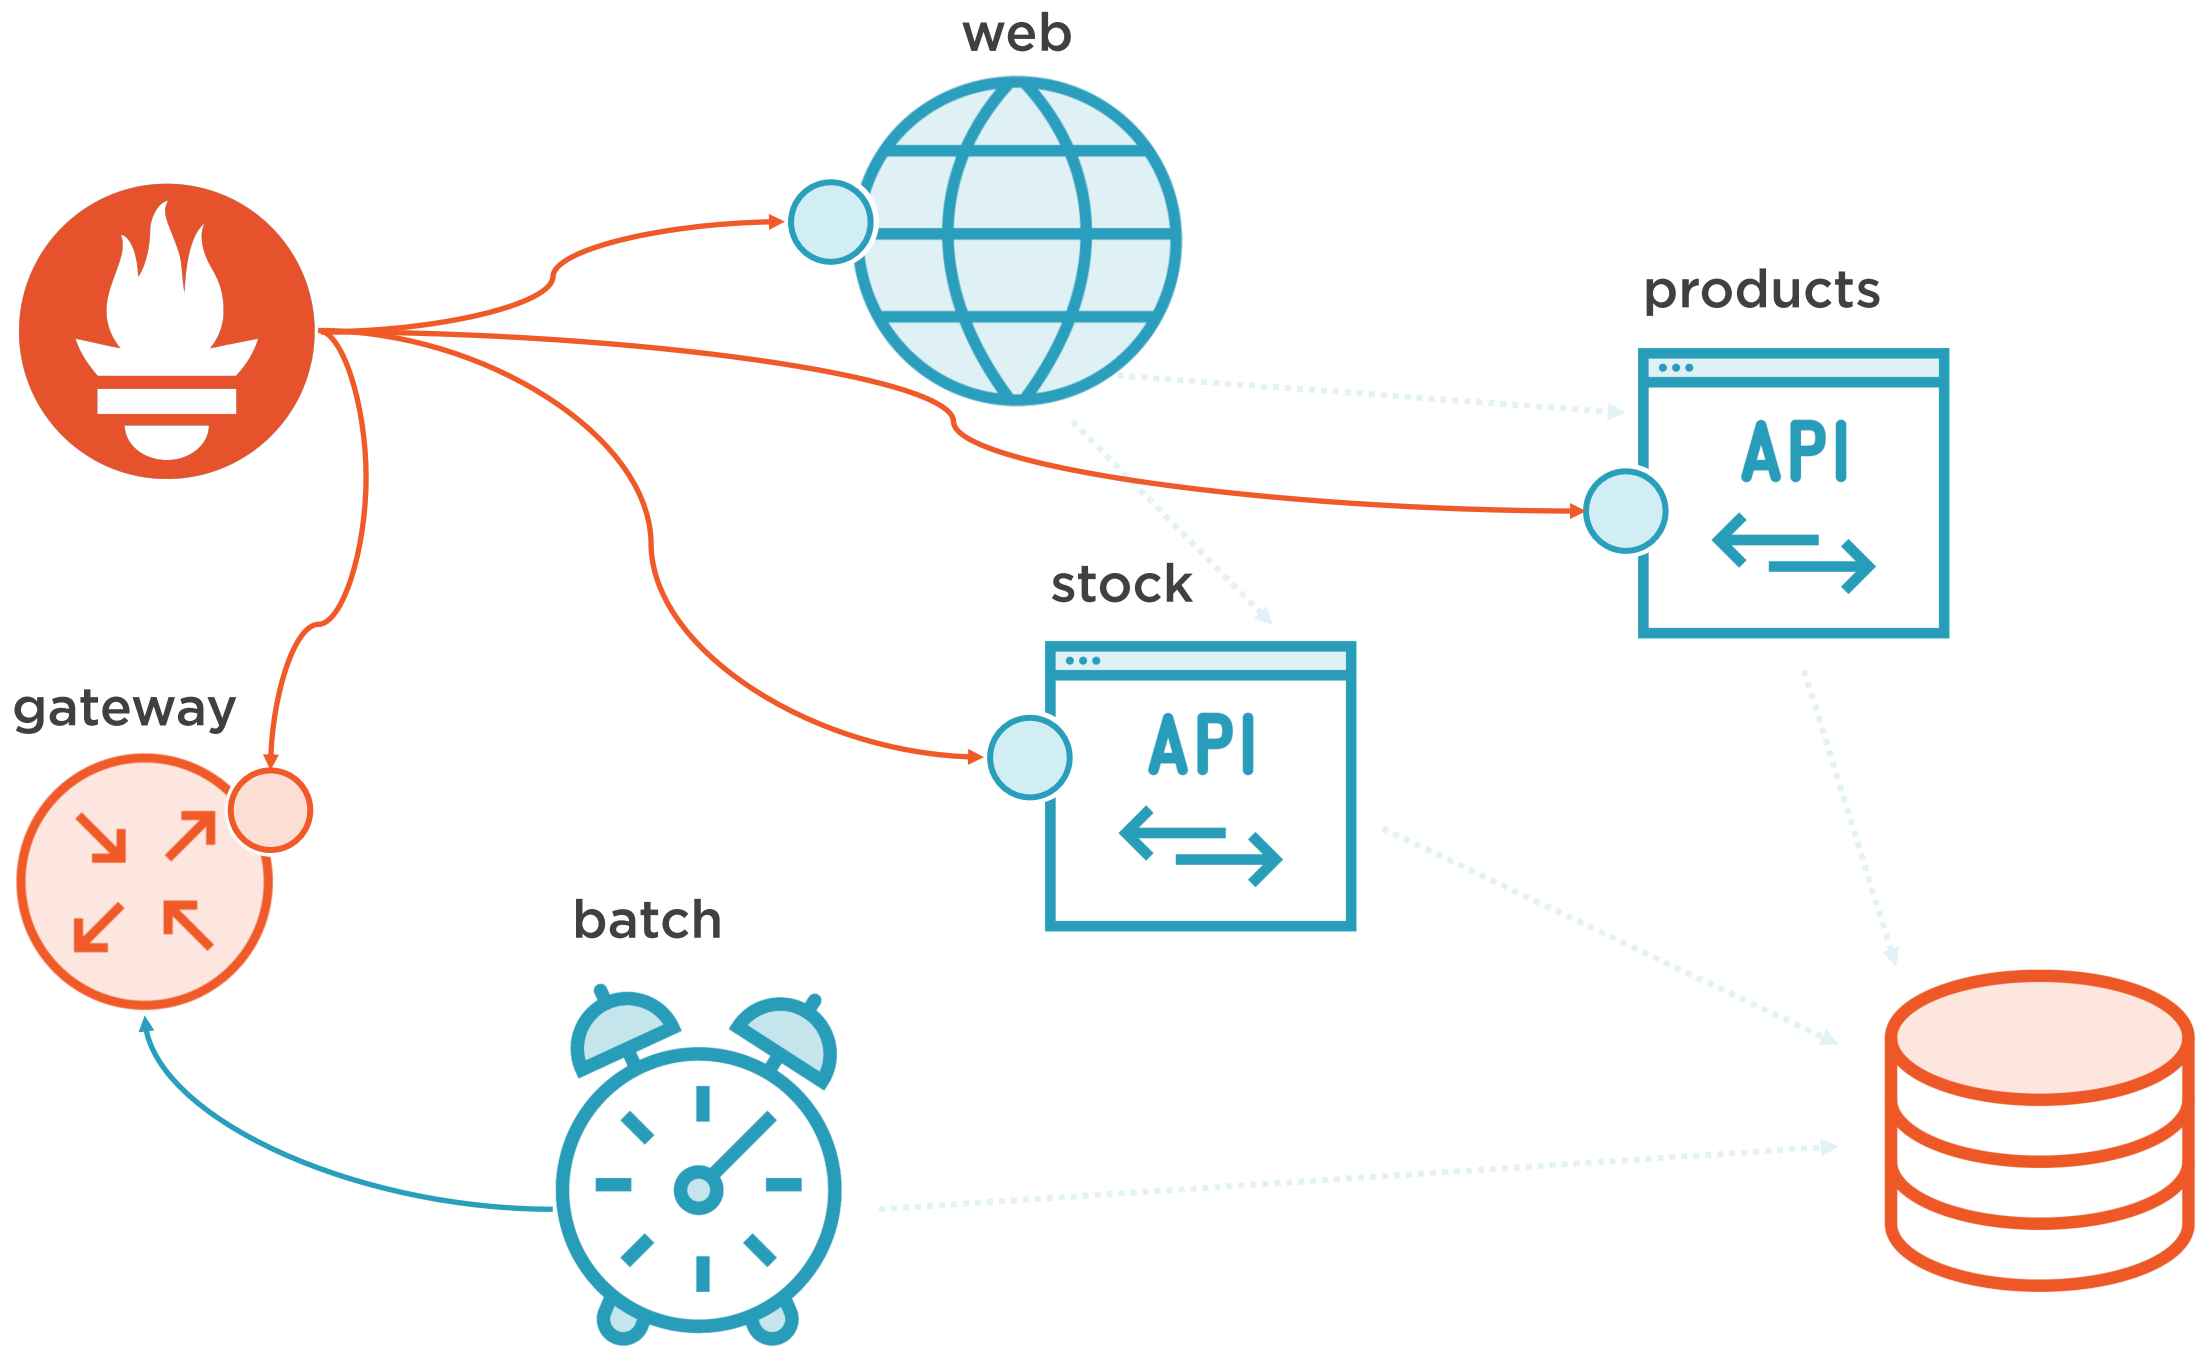












# Demo



## **Adding metrics to a web app**

- Using the .NET client library
- Wiring the metrics endpoint
- Verifying the metrics



## Library

Package reference



## Wiring

Plug into app runtime



## Metrics

Record custom values

# Reference the client library

WiredBrain.Web.csproj

```
<PackageReference  
  Include="prometheus-net.AspNetCore"  
  Version="3.6.0"/>
```

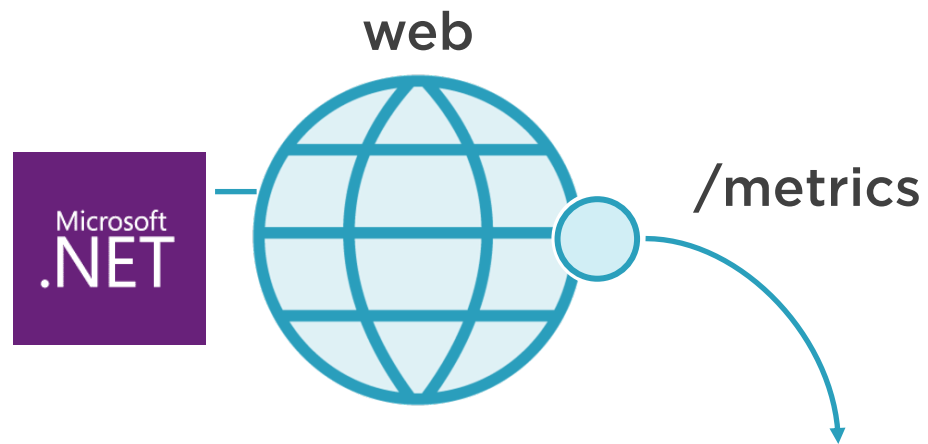
# Wire up the metrics

Startup.cs

```
// after app.UseRouting();
```

```
app.UseMetricServer();
```

```
app.UseHttpMetrics();
```



```
# HELP process_cpu_seconds_total Total user and system CPU time
```

```
# TYPE process_cpu_seconds_total counter
```

```
process_cpu_seconds_total 1.9
```

```
# HELP dotnet_total_memory_bytes Total known allocated memory
```

```
# TYPE dotnet_total_memory_bytes gauge
```

```
dotnet_total_memory_bytes 11622608
```

```
# HELP http_request_duration_seconds The duration of HTTP requests
```

```
# TYPE http_request_duration_seconds histogram
```

```
http_request_duration_seconds_bucket{code="200",method="GET",le="0.064"} 3
```

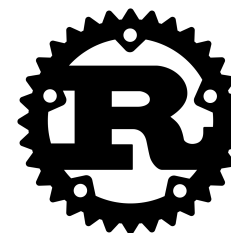


## Client libraries

Official



Community



...

<https://is.gd/awoded>





Client Library

Collector Registry



```
new Gauge("name", "text")
```



Collector



Collector



Collector



**Counter**  
Must be  
supported



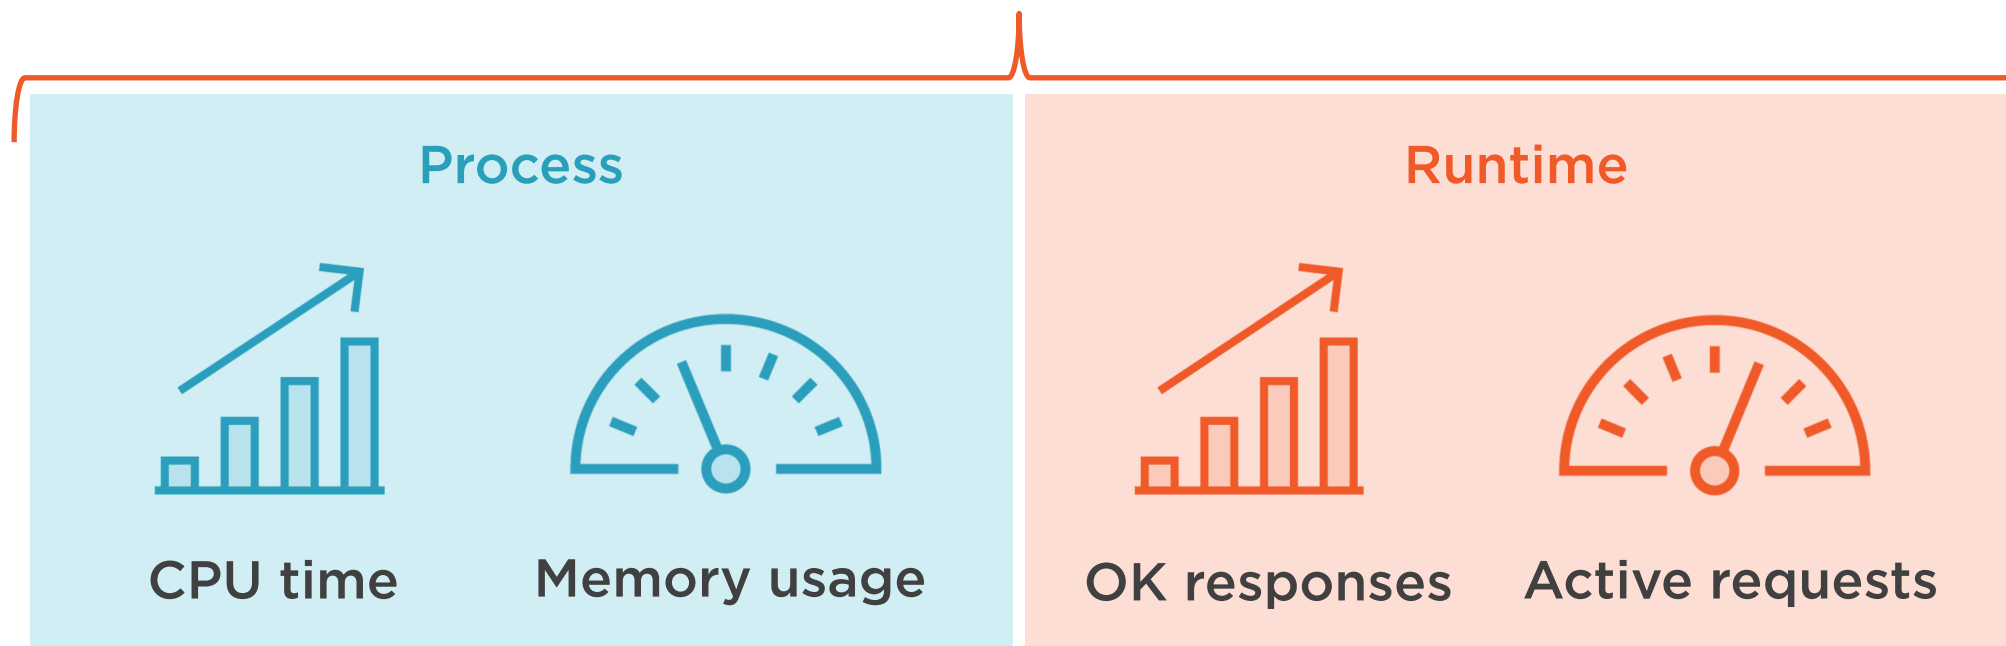
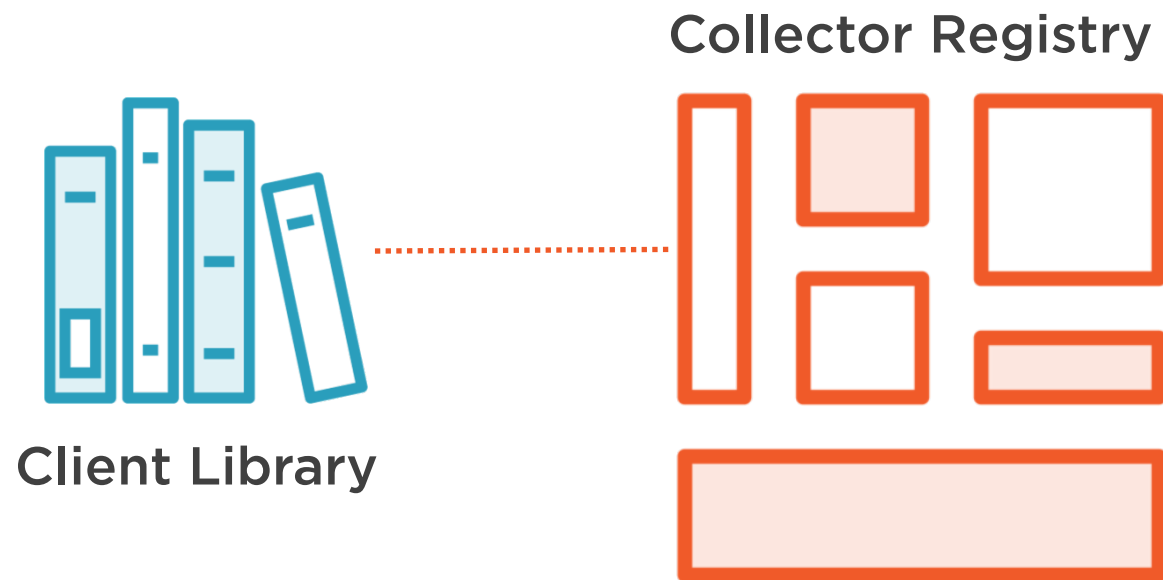
**Gauge**  
Must be  
supported



**Summary**  
Should be  
supported



**Histogram**  
Should be  
supported

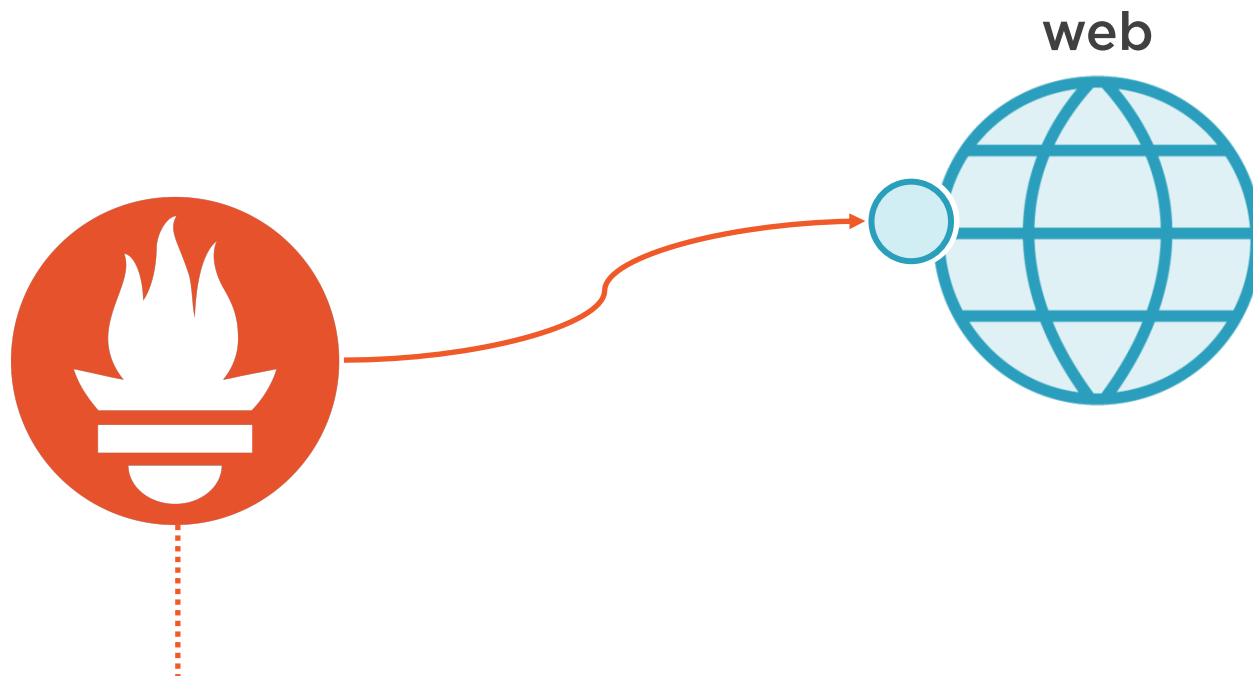


# Demo



## Exploring client library metrics

- Scrape the app with Prometheus
- Build simple graphs
- Run a load test



**http\_requests\_in\_progress**{method="GET",controller="Home",action="Index"} 15

**dotnet\_total\_memory\_bytes** 11622608

**http\_request\_duration\_seconds\_bucket**

{code="200",method="GET",controller="Home",action="Index",le="0.016"} 0

**http\_request\_duration\_seconds\_bucket**

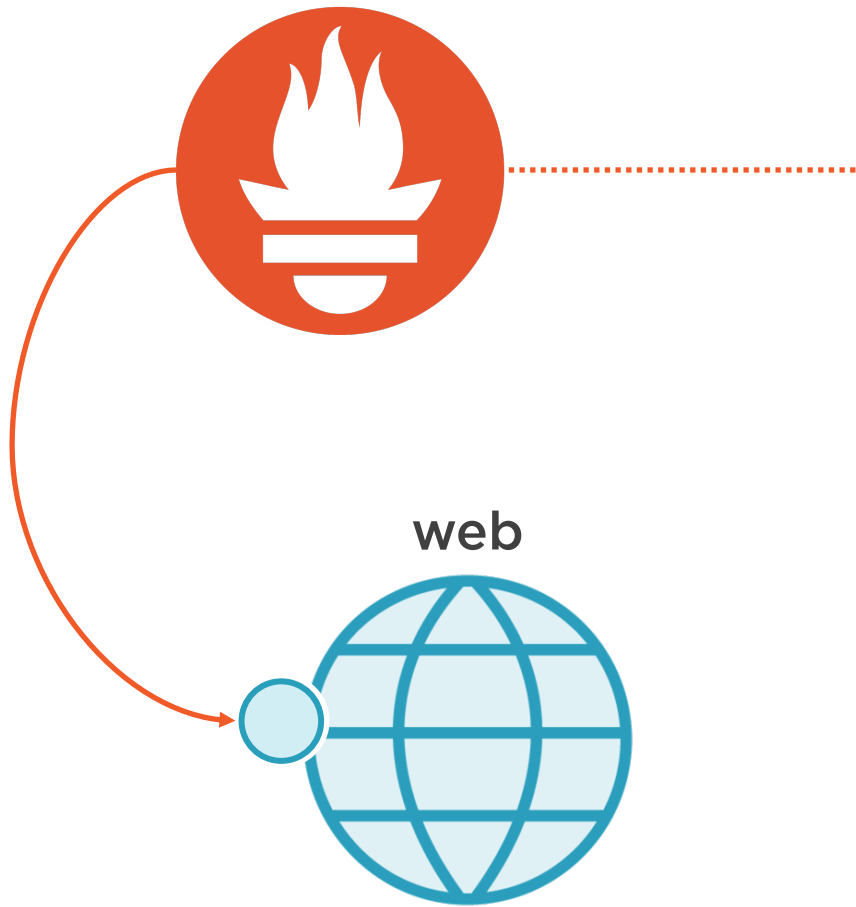
{code="200",method="GET",controller="Home",action="Index",le="0.032"} 2

# Wire up the metrics

Startup.cs

*// the client library collects HTTP metrics*

```
app.UseHttpMetrics();
```



**process\_cpu\_seconds\_total** 1.9

**process\_start\_time\_seconds** 1599471459.24

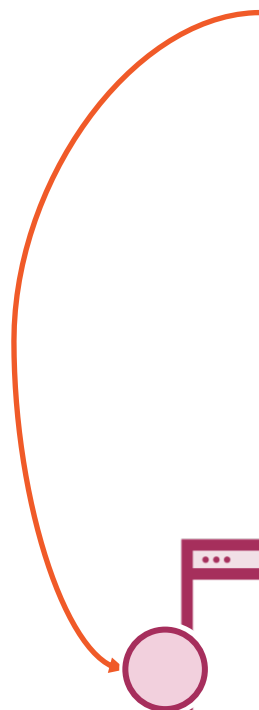
**process\_num\_threads** 26

**process\_open\_handles** 180

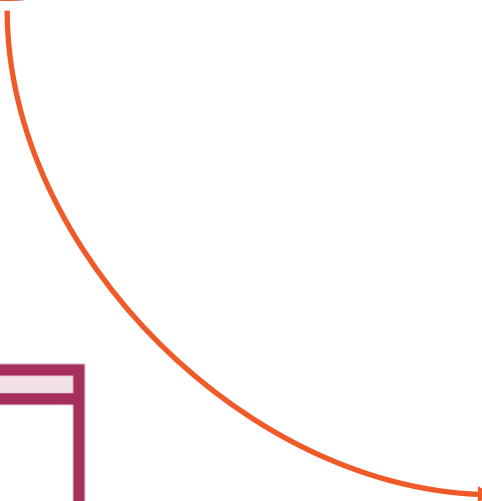
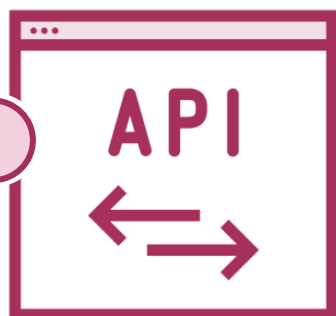
**dotnet\_collection\_count\_total{gen="0"}** 6

**dotnet\_collection\_count\_total{gen="1"}** 4

**dotnet\_collection\_count\_total{gen="2"}** 1



api



web



**dotnet\_total\_memory\_bytes**

{job="web",instance="w1"} 11622608

**dotnet\_total\_memory\_bytes**

{job="api",instance="a1"} 43102213





Client Library

## Collector Registry



### Process



Memory usage

### Runtime



Response time



Active requests

### Custom






# Site Reliability Engineering (SRE): The Big Picture

★★★★★ By Elton Stoneman

Site Reliability Engineering (SRE) is how Google runs production systems, promoting high availability with high velocity and removing operational toil. It achieves the same goals as DevOps without the culture shift.



## Course info

Rating	★★★★★ (29)
Level	Beginner 
Updated	Mar 5, 2020 
Duration	1h 41m 

## Description

Site Reliability Engineering (SRE) is a set of principles and practices that supports software delivery - keeping production systems stable and still delivering new features at speed. In this course, Site Reliability Engineering (SRE): The Big Picture, you'll get a thorough overview of how SRE works and why it's a good choice for many organisations. First, you'll learn the differences between SRE, DevOps, and traditional operations. Next, you'll discover how engineering practices help to reduce toil and provide more time to focus on high value tasks. Finally, you'll learn how SRE approaches monitoring and alerting, and about the SRE approach to managing incidents. When you're finished with this course, you'll be able to evaluate SRE and see if it's a good fit for your organisation.

## Service Levels, Monitoring, and Alerting

- 🔒 Understanding Service Level Objectives and Error Budgets
- 🔒 Defining Service Level Indicators and Service Level Objectives
- 🔒 Alerting on Service Level Objectives
- 🔒 Module Summary and SLO Improvement
- 🔒 Monitoring Service Level Indicators

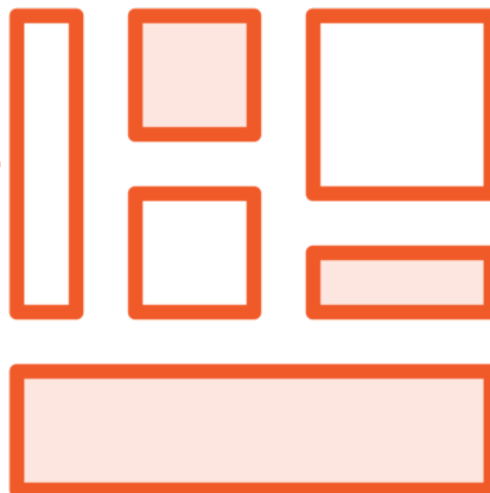
## Incident Management: On-call and Postmortems

<https://is.gd/veroto>



Client Library

## Collector Registry



`web_info {version="0.1.0"} 1`

### Process



Memory usage

### Runtime



Response time



Active requests

### Custom

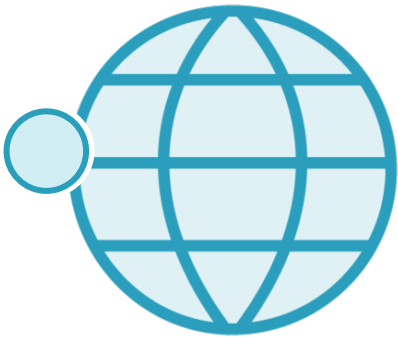


App info



```
dotnet_total_memory_bytes  
{version="0.1.0"} 28413694
```

web



```
dotnet_total_memory_bytes  
{version="0.2.0"} 10028317
```

# Demo

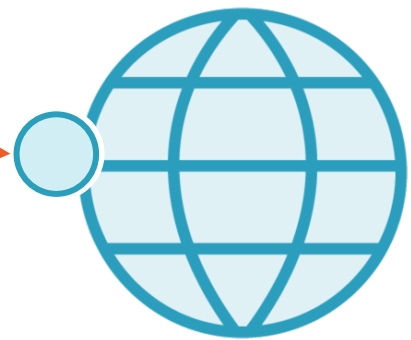


## Recording application information

- Adding a custom metric
- Using labels for version info
- Joining the info metric in queries



web



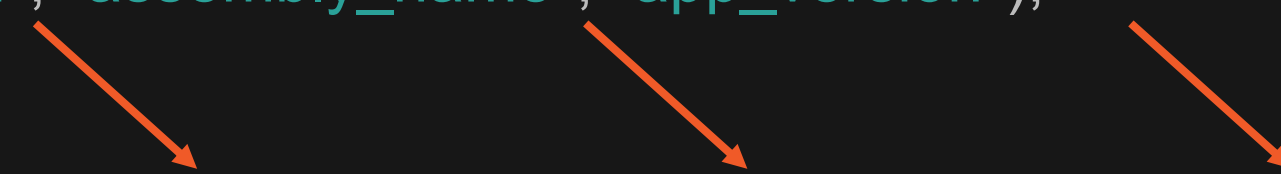
### **web\_info**

```
{dotnet_version="3.1.7",  
  assembly_name="WiredBrain.Web",  
  app_version="0.1.0"} 1
```

# Recording info metrics

Program.cs

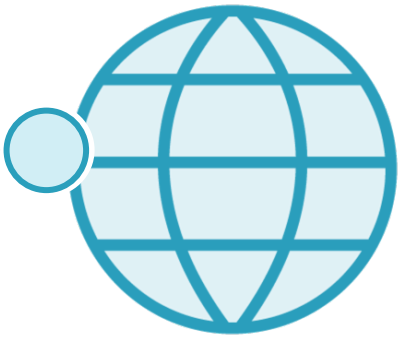
```
private static readonly Gauge _InfoGauge =  
    Metrics.CreateGauge("web_info", "Web app info",  
        "dotnet_version", "assembly_name", "app_version");  
//...  
_InfoGauge.Labels("3.1.7", "WiredBrain.Web", "0.1.0").Set(1);
```



The diagram consists of three orange arrows pointing from the parameters of the `CreateGauge` method call to the parameters of the `Labels` method call. The first arrow points from `"dotnet_version"` to `"3.1.7"`. The second arrow points from `"assembly_name"` to `"WiredBrain.Web"`. The third arrow points from `"app_version"` to `"0.1.0"`.



web



**http\_requests\_in\_progress 10**

**http\_requests\_in\_progress 150 @ 5s**

**http\_requests\_in\_progress 300 @ 10s**

**http\_requests\_in\_progress 900 @ 15s**

**http\_requests\_in\_progress 10 @ 20s**

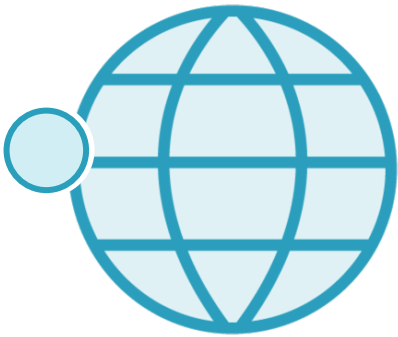




**http\_requests\_in\_progress**

{app\_version="0.1.0"} 10

web



**http\_requests\_in\_progress** 10

**web\_info**

{dotnet\_version="3.1.7",  
assembly\_name="WiredBrain.Web",  
app\_version="0.1.0"} 1

# Summary



## Instrumenting with client libraries

- Package reference
- Endpoint publishing
- Default metrics collection

## Custom metrics

- Standard Prometheus types
- Labels

## Info metrics

- Record application versions
- Surface in PromQL results

Up Next:

Recording custom application metrics

---